

EventTrace: Verifying Ordered Semantic Events for Reinforcement Fine-Tuning of Vision-Language Navigation

Anonymous CVPR submission

Paper ID *****

Abstract

Reinforcement fine-tuning of vision-language navigation (VLN) policies commonly relies on destination success or similarity to a reference route. These signals leave an instruction’s intermediate requirements unresolved: reaching the endpoint does not establish that the agent entered the requested room or passed the specified landmark. We study an online reward that represents an instruction as an ordered list of semantic events and verifies each event against observations before and after executed actions. A frozen vision-language model supplies a three-way verdict (complete, incomplete, uncertain); motion and stopping checks reject visually plausible but unexecuted events. Each event is rewarded at most once, and the process reward is bounded below the success bonus. We integrate EventTrace into multi-turn ActiveVLN/GRPO and compare it with destination-only training for 64 optimization steps across three matched seeds. All six checkpoints and the SFT initialization are evaluated on the same 1,839 R2R val-unseen episodes. EventTrace changes success rate by +0.65, +0.44, and −2.88 percentage points across seeds; its mean paired change is −0.60 points (sample standard deviation 1.98). A single-AI blind review of 49 selected transitions, replayed with training-time image encoding, finds that 7 of 15 verifier completion predictions match its labels. Two subsequent independent human reviewers agree on only 33 of these 49 cases; their 16 disagreements remain adjudicated. The method is operational, but these results do not establish a navigation improvement or reliable semantic verification.

1. Introduction

Vision-language navigation asks an embodied agent to follow a natural-language route through visual observations. The Room-to-Room (R2R) task introduced

instruction-guided navigation in real indoor environments [2]; VLN-CE moves this task into continuous control [7]. Success rate (SR) and success weighted by path length (SPL) capture destination attainment, yet a successful stop can coexist with skipped intermediate instructions. Instruction fidelity work made this gap explicit [6].

Recent reinforcement fine-tuning (RFT) approaches give visual-language policies online or trajectory-level feedback. ActiveVLN performs multi-turn exploration with group-relative policy optimization (GRPO) [14, 17]. VLN-R1 uses time-decayed action rewards [11]; Nav-R1 combines format, understanding, and navigation rewards [9]; SeeNav-Agent proposes step-level verifiable process rewards in another navigation benchmark [16]. Progress-Think learns semantic progress from visual histories and instruction prefixes [15]; SACA audits stepwise landmark progress and routes mixed-success and all-failure groups differently [8]. Our question is narrower: can a navigation instruction’s ordered, physical events provide a bounded reward during closed-loop R2R-CE rollouts, without matching the policy’s actions to a reference path?

We introduce EventTrace, a modular reward on top of an existing multi-turn policy. The reward uses a short event list, a frozen visual verifier with an abstention state, deterministic checks of motion or stopping, and a once-only order-preserving tracker. The actor sees only the original egocentric observations and instruction; simulator displacement is available only to the reward. Figure 1 illustrates the distinction between seeing a bathroom and actually entering it.

We evaluate the reward with matched three-seed training and complete R2R val-unseen coverage. The observed seed variation and an exploratory verifier audit limit the empirical claim: this implementation is reproducible but does not improve navigation consistently.




Figure 1. Real simulator frames from a diagnostic replay of R2R training episode 2889. Visibility alone does not complete the event “enter bathroom.” The verifier uses before/after views and executed motion; the ordered tracker later checks “stop near sink.” These frames illustrate the mechanism and are not an independent accuracy test.

2. Related Work

Path and instruction fidelity. SR and SPL reward the destination and efficiency, whereas nDTW and coverage weighted by length score (CLS) compare trajectories with a reference route [5, 6]. Reference trajectories can be useful supervision but do not directly answer whether a distinct visual event actually occurred when multiple valid routes exist. Our reward targets event completion during execution; route metrics remain complementary evaluation measures.

RL for VLN. ActiveVLN is our implementation base because its environment can return feedback after each multi-action turn [17]. VLN-R1 and Nav-R1 show other ways to shape VLN rewards [9, 11]. Progress-Think aligns observation and instruction progress and co-trains a progress-aware policy [15]; SACA combines a landmark auditor with scenario-dependent group optimization [8]. SeeNav-Agent emphasizes step-level verification [16]. EventTrace tests a narrower combination of a once-only ordered event state, an explicit uncertainty outcome, and motion gates against rewarding a landmark merely seen without the requested transition. Our pilot applies GRPO to trajectory-level returns; it does not claim turn-level credit assignment or stronger navigation performance than these systems.

3. Method

3.1. Task and event representation

An instruction x yields a list $E(x) = (e_1, \dots, e_K)$, where e_k contains a verb type, target phrase, and source span. Supported event types include enter, pass, descend, and stop_near. A frozen vision-language model (VLM) proposes at most three events, restricted to source spans in the instruction. The eight-episode

Table 1. Motion gates in the pilot. The VLM must also verify that the image pair supports the specified target.

Event	Count	Required gate
enter	3	Forward action; $\ \Delta p\ \geq 0.10$ m
pass	3	Forward action; $\ \Delta p\ \geq 0.10$ m
descend	3	Vertical change $\Delta y \leq -0.08$ m
stop_near	8	Executed STOP action

wiring pilot uses curated lists; the main 64-step training runs parse instructions automatically at episode start. Parser errors produce an empty list and no semantic reward. We have not independently audited parser accuracy.

In the curated pilot, each of the eight episodes contains a stopping event; there are also three enter, three pass, and three descend events, for 17 in total. These counts describe the pilot distribution, not the prevalence of event types in R2R. Table 1 shows the deterministic portion of the verifier. Each gate can reject completion before a VLM call, but passing a gate is never sufficient to award a point.

At turn t , the policy receives the instruction, action history, and current egocentric image I_t , and emits up to three low-level actions. The simulator executes these actions and returns I_{t+1} , pose change Δp_t , and task metrics. Only the current pending event e_{k_t} is queried. The frozen verifier produces

$$v_t = V(e_{k_t}, I_t, I_{t+1}, a_t, \Delta p_t) \in \{Y, N, U\}, \quad (1)$$

where U denotes insufficient evidence. A view of a landmark is not by itself a completion verdict. For crossing and passing, a forward displacement check is required; descending requires a negative vertical displacement; stopping near a target requires an executed STOP. The VLM inspects whether the image pair supports the specified target and transition. These are deliberately conservative gates, not a learned proof of physical state.

3.2. Ordered, bounded process reward

Let $k_t \in \{1, \dots, K+1\}$ index the first unfinished event. The event fires only when its verifier says Y. Then

$$c_t = \mathbb{I}[k_t \leq K \wedge v_t = Y], \quad (2)$$

$$k_{t+1} = k_t + c_t, \quad r_t^{\text{sem}} = \frac{\lambda}{K} c_t. \quad (3)$$

A N or U verdict leaves k_t unchanged. The tracker prevents repeated reward for the same landmark, enforces event order, and guarantees $\sum_t r_t^{\text{sem}} \leq \lambda$. Because only the pending event is checked, at most one event fires per turn even if the agent performs several

actions. This avoids ambiguous credit when a model emits an action batch but can delay recognition of adjacent events. The tracker retains a confirmed event after later image changes, so a completed event cannot be unconfirmed by viewpoint drift.

The base destination term is the weighted-success reward already used by ActiveVLN, with nDTW weight zero. We add a fixed success floor b to both arms:

$$R(\tau) = \sum_t r_t^{\text{sem}} + S(\tau) \left[b + \alpha \left(1 - \min \left(\frac{d_T}{3}, 1 \right) \right) \right], \quad (4)$$

where $S(\tau)$ is simulator success after STOP, d_T is terminal distance in meters, $\alpha = 15$, $b = 2$, and $\lambda \in \{0, 1\}$. Hence any success receives at least 2 while a failure receives at most 1 semantic point. This priority holds for this nonnegative reward configuration; it does not imply that the verifier is correct. The bound also makes failure reward interpretable: partial event completion can distinguish two unsuccessful attempts without allowing a failed route to outrank a successful one solely through event bonuses. The STOP event is distinct from simulator success: it demands visual proximity to its instruction target, while $S(\tau)$ uses the simulator’s goal region. Disagreement between these signals is precisely a case the proposed reward aims to expose.

For each training episode, we sample two trajectories and normalize their returns within the group for GRPO [14]. Event rewards are emitted after each turn, but the current trainer sums them before constructing a trajectory advantage. Consequently, the implementation changes which trajectories are preferred, not which individual turn receives the policy gradient. Turn-level event timing is available in the log for future credit-assignment work.

3.3. Verifier failure policy

The verifier is a frozen Qwen3-VL-8B-Instruct model. A service timeout, malformed response, or uncertain decision yields zero event reward; it does not end the episode. The reward requires no reference action prefix. Simulator pose is used only for reward verification, never added to the actor observation. In the pilot, the fixed event manifest is matched to both episode ID and exact instruction text to avoid assigning another episode’s event list; the main runs use online parsing. The verifier sees two RGB frames resized to fit within 448×448 pixels, the event type and target, its source phrase, the executed action text, translation magnitude, and vertical displacement. Its output is a deterministic one-letter Y/N/U response; generated explanations are not treated as confidence scores. This makes the reward externally auditable through a saved per-turn verdict, although a forced-choice label is still

fallible. The HTTP verifier is separated from the simulator so it can be frozen and logged independently of policy updates.

4. Experimental Setup

Environment and policy. We use R2R train in VLN-CE, rendered from Matterport3D scans [3], the ActiveVLN implementation, and an SFT-initialized Qwen2.5-VL-3B policy. Matterport3D data and derived imagery are subject to its [academic terms of use](#). Each turn contains at most three grouped action commands. The configured limits are 12 turns and 36 commands; the strict budget check can permit one extra command, and a command may execute several simulator steps. Actor sampling uses temperature 1.2 and two independently sampled rollouts per item. GRPO trains on two A800 GPUs at learning rate 10^{-6} with batches of four episodes. The destination-only control sets $\lambda = 0$; the event arm sets $\lambda = 1$. Both arms have the same SFT initialization, destination reward, data order within seed, and optimizer budget.

Pilot and corrected main training. We first replay eight training episodes with expert actions solely to inspect expected events, then run a two-step reward-wiring pilot with curated events. These replays are neither held-out labels nor training results for the main comparison. For the main study, each arm runs 64 optimization steps at seeds 11, 22, and 33: 256 train-episode exposures and 512 sampled rollouts per arm and seed. The event arm uses automatically parsed events. A preflight check detected that a fixed per-request sampling seed duplicated every two-rollout GRPO group and yielded zero actor gradients; we excluded that attempt, removed the per-request seed while retaining the engine seed, and reran all six models. Every corrected group has distinct sampled actions. The archive records the excluded attempt and corrected patch.

Full val-unseen evaluation. SR is the fraction of trajectories satisfying the simulator goal condition; SPL also accounts for path efficiency. Mean terminal goal distance is in meters. Training rollout success and return are diagnostics, not held-out outcomes; returns differ in definition between arms. The SFT model and six corrected step-64 checkpoints run on the identical manifest of all 1,839 R2R val-unseen episodes from 11 scenes. We use one generation per episode, temperature 0.2, top- p 0.8, a 12-turn limit, and at most 76,800 image pixels. Four Habitat shards per checkpoint divide the manifest without overlap. The evaluator uses

Table 2. Observed pilot results on eight training episodes. Each arm ran two optimizer steps and produced 16 sampled rollouts. The event arm’s larger mean return is expected from the added reward and does not indicate better navigation.

Metric	Destination only	EventTrace
Successful rollouts	13/16	13/16
Mean trajectory return	9.325	10.455
Mean event progress	–	0.656
Confirmed event hits	–	23
Uncertain verifier turns	–	33
Verifier service errors	–	0

the same prompt and environment-specific OpenAI-compatible request and image-resize shim for all seven models. The completion gate checks the exact episode-ID set and zero inference errors. Episode IDs are paired within each seed, but stochastic decoding uses only one generation per checkpoint and is an additional source of variation.

Blind verifier review. Separately, we collect shortest-path follower trajectories from 16 different val-unseen episodes, of which 12 yield at least one automatically parsed event. Before querying the verifier, one AI-assisted reviewer inspects RGB image pairs, executed actions, and motion and fixes 49 three-way completion labels across those 12 episodes. For the reported replay, both frames use the training client’s 448-pixel JPEG encoding. The sample includes apparent transitions, controls, and final stops. The original AI review was independent of verifier predictions but had no human ground truth. Subsequently, two independent humans labeled the same answer-free, randomized 49-item package without access to predictions or prior labels. Their files were locked before comparison. We report their agreement separately and do not treat unresolved labels as verifier ground truth. Selected class proportions are not representative of deployment traffic.

5. Results and Analysis

Reward reaches optimization. In the first optimizer step, both arms sample the same eight action sequences and attain seven successes. For every matched trajectory, the difference in total return equals its semantic event score, including a failed episode 7110 trajectory that receives 0.5 for completing “descend stairs” but never reaches the destination. The modified return is consumed by the training reward manager, and both arms save step-2 checkpoints after nonzero policy updates. This verifies the wiring of Eq. 4; it is not evi-

dence that event reward improves the policy.

What the verifier recognizes. In diagnostic expert replays, four of eight episodes receive full event credit, two receive partial credit, and two receive none; the tracker confirms 11 of 17 curated events. For episode 2889, seeing the sink from outside the bathroom produces uncertain verdicts before entrance, while later motion and view changes confirm enter. A rotation followed by premature STOP yields zero semantic credit in a negative replay. There are also false negatives: one shallow stair descent is rejected by the vertical-motion threshold, and one landmark is never confirmed. Because the events were curated using these episodes and the replays lack independent annotations, 11/17 must not be reported as verifier recall or accuracy.

Main training diagnostics. All six corrected runs finish 64 steps, with 256 diverse GRPO groups per arm and seed. Training-rollout success is 111/512 versus 119/512 (control versus event) for seed 11, 108/512 versus 119/512 for seed 22, and 88/512 versus 82/512 for seed 33. These on-policy samples change as training proceeds, so they are not comparable to the held-out evaluation. The event verifier records 272, 275, and 243 confirmed event hits over the three event-arm runs. These counts are not event accuracy: events are automatically parsed and no independent labels cover these rollouts.

Complete val-unseen results. Table 3 gives all seven checkpoints on the same 1,839 episodes. Within matched seeds, EventTrace changes SR by +0.65, +0.44, and –2.88 percentage points, and SPL by +0.58, +0.49, and –3.49 points. The mean paired seed change is -0.60 ± 1.98 points for SR and -0.80 ± 2.32 points for SPL (mean $\pm$ sample standard deviation over three seeds). At the episode level, event-only versus control-only successes are 157 versus 145, 197 versus 189, and 155 versus 208. This variation does not support a consistent improvement. All seven evaluations have exact 1,839-episode coverage and zero logged inference errors. The earlier 16-episode, two-step probe yielded 6/16 successes for every arm; it was exploratory and is superseded by the complete comparison.

Verifier failure analysis. Table 4 compares the frozen verifier with labels fixed before its predictions were obtained. Seven of eight reviewed completions receive Y, but only 7 of 15 Y predictions match a reviewed completion; four are on clear non-completions and four on ambiguous turns. One false positive occurs after the

Table 3. Complete R2R val-unseen evaluation. Every row uses the same 1,839 episodes from 11 scenes, one generation per episode, and zero inference errors. Goal distance is in meters. C: destination-only control; E: EventTrace.

Checkpoint	Successes	SR (%)	SPL (%)	Dist.
SFT	544	29.58	28.69	7.03
Seed 11 C	479	26.05	25.32	7.75
Seed 11 E	491	26.70	25.91	7.63
Seed 22 C	566	30.78	30.14	6.91
Seed 22 E	574	31.21	30.63	6.98
Seed 33 C	502	27.30	26.65	7.14
Seed 33 E	449	24.42	23.16	7.76

Table 4. Training-encoded replay of 49 selected val-unseen event transitions from 12 episodes. Rows are blind review labels and columns are frozen verifier predictions. U means evidence insufficient; the single-AI review has no independent human adjudication.

	Pred. Y	Pred. N	Pred. U
Review Y	7	0	1
Review N	4	13	12
Review U	4	1	7

agent has already entered the requested room, exposing a failure to distinguish current state from an event newly completed this turn. Motion gates alone do not eliminate this error. These selected examples cannot establish deployment precision, but they make verifier reliability a concrete limitation of the present reward.

Independent human review. The two blinded human reviewers agree on 33/49 three-way labels (67.35%; Cohen’s $\kappa = 0.459$). Agreement is 17/23 for enter, 6/10 for pass, and 10/16 for stop_near; 16 cases remain disputed. The brief rationales include interpretations of motion toward a landmark as completion and cases where a previously visited landmark cannot be established from the supplied before/after views. This shows that the local evidence and event rubric can be under-determined; it does not isolate model error from label ambiguity. No model-versus-human precision or recall is computed before independent adjudication, and this deliberately selected pilot cannot estimate population accuracy.

6. Discussion and Limitations

The main risk is reward misspecification: the VLM may confuse doorways, and displacement gates may miss small valid movements. One-at-a-time verification may miss events completed together. Before/after frames limit timing to a turn; trajectory-level GRPO

Table 5. Completed full val-unseen comparisons, each on 1,839 episodes. Entries are paired mean $\pm$ sample SD over three training seeds, in percentage points. Each row has its own matched control. EventTrace uses 64 steps and two rollouts per group; the follow-ups use 128 steps, with two rollouts for policy-prefix training and four for the last three rows. The group-size comparison uses a two-rollout control and twice its rollout compute. Differences between rows cannot rank methods.

Comparison	Δ SR	Δ SPL
EventTrace	-0.60 ± 1.98	-0.80 ± 2.32
Policy-prefix	-1.14 ± 5.41	-0.55 ± 5.75
Group four versus two	$+2.10 \pm 1.93$	$+2.51 \pm 2.27$
Privileged turn-wise	$+2.47 \pm 3.18$	$+2.51 \pm 2.78$
Terminal-RLOO	$+1.61 \pm 4.22$	$+1.07 \pm 3.60$

offers limited credit assignment to the decisive action.

The complete val-unseen comparison shows a negative mean effect and large seed variation, but 64 steps expose only 256 training-episode instances per run. For EventTrace, a larger training budget and repeated decoding have not been evaluated. The selected 49-example review now has two independent human labelers, but their disagreements have not been adjudicated; instruction parsing is also unaudited. We cannot isolate noisy verification, parser errors, reward scale, budget, or optimization as the cause. Separately sampled and adjudicated verifier and parser sets, followed by ablations of ordering, abstention, gates, reward budget, and turn-level advantage, are needed. This audit cannot establish event precision or recall.

6.1. Evidence from subsequent training changes

The follow-ups change credit assignment, training-state distribution, rollout group size, or the reward representation. They therefore have separate matched controls and development protocols; they are not ablations that identify why EventTrace failed. Table 5 summarizes the completed three-seed comparisons. Increasing rollout count gives a positive mean with additional simulator work, whereas policy-prefix training has a negative mean. Privileged turn-wise credit and normalized terminal-RLOO have positive means but opposite seed effects and exploratory intervals spanning zero. None validates an observation-only semantic reward. Detailed protocols, unsuccessful screens, and individual seed results are retained in the supplementary material.

The one-seed PPO/GAE implementation also regressed on val-seen. Its recurrence discounts generated action-text tokens, so this result does not evaluate decision-clock GAE or establish that actor-critic methods are unsuitable for VLN. Observation-only reward

probes improved some offline rankings but failed on-line navigation or conservative calibration gates. These outcomes motivate separating evidence quality from the policy update rule rather than increasing group size alone.

6.2. Positive trajectory weighting and pending replication

Positive-only trajectory update. After the inconsistent terminal-RLOO replication, we froze a different on-policy update on 45 R2R-train fit scenes, with eight development and eight reserved scenes. Both arms receive the same distance-weighted success-15 plus reference-route nDTW-5 terminal score R_i . For four sampled trajectories, the candidate uses

$$Q_i = R_i + 15S_i,$$

$$A_i = M_i \min \left\{ 1.5, \frac{[Q_i - \frac{1}{3} \sum_{j \neq i} Q_j]_+}{5} \right\}, \quad (5)$$

where S_i is simulator success and M_i is one only for successful trajectories or failures with $R_i \geq 2.5$. This prioritizes success and removes negative trajectory credit, inspired by self-imitation learning [10]. The update uses fresh on-policy groups and a peer-return baseline, without a replay buffer or learned value function. Each weight is broadcast to action tokens, with observation tokens masked. The matched GRPO control and candidate use the same SFT initialization, 512 fit rows, 64 steps, group size four, one policy epoch, sequence-mean loss aggregation, and reference KL coefficient 0.01. Every step has a nonzero actor gradient in both arms.

On the frozen 256-episode, eight-scene development screen, candidate/control successes are 51/17 (SR 19.92%/6.64%); paired SR and SPL both improve by 13.28 percentage points. Both arms have exact coverage and zero inference errors, and local independent recount agrees. A CPU navmesh reconstruction matches every raw SPL: all successful paths are shorter than the initial goal distance, so their SPL is one under the 3 m success region. Descriptive scene-bootstrap intervals are [8.71, 19.35] points for both differences. The first evaluation’s wrong nDTW reference split was corrected before this complete sweep; all partial outputs and failed logs were archived. The original joint +2-point gate passed, and new matched 512-row, 128-step, group-four runs at seeds 11, 22, and 33 have begun. Their reserved 256-episode evaluation has not started. This is a one-seed development gain relative to GRPO, with low absolute success; the SFT initialization has not been separately evaluated on this screen. Earlier SFT and exploratory work may have seen these train scenes. The result does not establish unseen-scene ben-

efit or validate a deployable semantic verifier; its reward uses privileged simulator success and a reference route.

The replication protocol additionally fixes all six step-128 checkpoints for complete val-unseen evaluation, regardless of reserved-screen results. The unchanged SFT initialization will be evaluated on development, reserved, and val-unseen episodes as a shared reference. This extra protocol was frozen after the pilot and before scale evaluation. A single SFT decode is not three independent SFT training seeds. The reserved scenes are disjoint from this fit but may have been exposed during earlier SFT or exploratory work; val-unseen has also been used adaptively. Pending replication and SFT comparisons supply no additional observed gain.

7. Conclusion

We formulated a path-free semantic process reward for closed-loop VLN that verifies ordered instruction events against executed motion and before/after views. A wiring pilot confirms that the bounded reward enters online GRPO. In three matched 64-step training runs and complete 1,839-episode val-unseen evaluation, EventTrace has a mean paired SR change of -0.60 percentage points relative to destination-only training, with inconsistent seed effects. Two independent human reviewers agree on 33/49 selected transitions, but the remaining cases lack adjudication and the selected sample cannot establish verifier accuracy. The current method is reproducible; reliable semantic verification and a navigation benefit remain unproven.

References

- [1] Arash Ahmadian, Chris Cremer, Matthias Gallé, Marzieh Fadaee, Julia Kreutzer, Olivier Pietquin, Ahmet Üstün, and Sara Hooker. Back to basics: Revisiting REINFORCE style optimization for learning from human feedback in LLMs, 2024. 8
- [2] Peter Anderson, Qi Wu, Damien Teney, Jake Bruce, Mark Johnson, Niko Sunderhauf, Ian Reid, Stephen Gould, and Anton van den Hengel. Vision-and-language navigation: Interpreting visually-grounded navigation instructions in real environments. In CVPR, 2018. 1
- [3] Angel Chang, Angela Dai, Thomas Funkhouser, Maciej Halber, Matthias Niessner, Manolis Savva, Shuran Song, Andy Zeng, and Yinda Zhang. Matterport3D: Learning from RGB-D data in indoor environments. In International Conference on 3D Vision (3DV), 2017. 3
- [4] Jian Hu, Jason Klein Liu, Haotian Xu, and Wei Shen. REINFORCE++: Stabilizing critic-free policy optimization with global advantage normalization, 2025. 8

- [5] Gabriel Ilharco, Vihan Jain, Alexander Ku, Eugene Ie, and Jason Baldridge. General evaluation for instruction conditioned navigation using dynamic time warping. In NeurIPS, 2019. 2
- [6] Vihan Jain, Gabriel Magalhaes, Alexander Ku, Ashish Vaswani, Eugene Ie, and Jason Baldridge. Stay on the path: Instruction fidelity in vision-and-language navigation. In ACL, 2019. 1, 2
- [7] Jacob Krantz, Erik Wijmans, Arjun Majumdar, Dhruv Batra, and Stefan Lee. Beyond the nav-graph: Vision-and-language navigation in continuous environments. In ECCV, 2020. 1
- [8] Haoyuan Li, Rui Liu, Hehe Fan, and Yi Yang. Let’s reward step-by-step: Step-aware contrastive alignment for vision-language navigation in continuous environments, 2026. 1, 2
- [9] Qingxiang Liu, Ting Huang, Zeyu Zhang, and Hao Tang. Nav-r1: Reasoning and navigation in embodied scenes, 2025. 1, 2
- [10] Junhyuk Oh, Yijie Guo, Satinder Singh, and Honglak Lee. Self-imitation learning. In Proceedings of the 35th International Conference on Machine Learning, pages 3878–3887. PMLR, 2018. 6
- [11] Zhangyang Qi, Zhixiong Zhang, Yizhou Yu, Jiaqi Wang, and Hengshuang Zhao. Vln-r1: Vision-language navigation via reinforcement fine-tuning, 2025. 1, 2
- [12] John Schulman, Philipp Moritz, Sergey Levine, Michael I. Jordan, and Pieter Abbeel. High-dimensional continuous control using generalized advantage estimation. In International Conference on Learning Representations, 2016. 9
- [13] John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms, 2017. 9
- [14] Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, et al. Deepseekmath: Pushing the limits of mathematical reasoning in open language models, 2024. 1, 3
- [15] Shuo Wang, Yucheng Wang, Guoxin Lian, Yongcai Wang, Maiyue Chen, Kaihui Wang, Bo Zhang, Zhizhong Su, Yutian Zhou, Wanting Li, Deying Li, and Zhaoxin Fan. Progress-think: Semantic progress reasoning for vision-language navigation. In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition, pages 4076–4086, 2026. 1, 2
- [16] Zhengcheng Wang, Zichuan Lin, Yijun Yang, Haobo Fu, and Deheng Ye. Seenav-agent: Enhancing vision-language navigation with visual prompt and step-level policy optimization, 2025. 1, 2
- [17] Zekai Zhang, Weiye Zhu, Hewei Pan, Xiangchen Wang, Rongtao Xu, Xing Sun, and Feng Zheng. Activevln: Towards active exploration via multi-turn rl in vision-and-language navigation, 2025. 1, 2

EventTrace: Verifying Ordered Semantic Events for Reinforcement Fine-Tuning of Vision-Language Navigation

Supplementary Material

546 A. Scope of the supplementary evidence

547 This material preserves the exploratory protocols,
548 failed screens, training audits, and post-result diagnos-
549 tics summarized in the main text. All reported nav-
550 igation differences have episode-matched controls and
551 explicit coverage checks. The screens differ in episode
552 sets, budgets, and prior use; their absolute success
553 counts cannot rank methods across experiments. No
554 unadjudicated semantic labels or offline reward rank-
555 ings are treated as navigation gains.

556 B. Credit-estimator and group-size develop- 557 ment

558 Turn-level credit assignment pilot. We tested a group-
559 four, actor-only return-to-go/leave-one-out estimator
560 inspired by RLOO [1] and REINFORCE++ [4]. Each
561 executed turn receives its remaining return rela-
562 tive to active peer rollouts; observation tokens are
563 masked. The candidate also adds training-only simu-
564 lator geodesic progress, so both credit assignment and
565 reward change. From the same SFT state and 256 train
566 rows, it and a destination-only GRPO control ran 64
567 steps; all candidate actor gradients were nonzero. On
568 a frozen 256-episode val-seen manifest (53 scenes, ex-
569 act coverage, zero inference errors), candidate/control
570 successes were 68/78: paired SR -3.91 and SPL -3.72
571 points. The prespecified joint $+2$ -point gate failed, so
572 this mechanism was not expanded. One seed and de-
573 code cannot isolate turn-level credit assignment or es-
574 tablish unseen-scene gains. This is not a deployable
575 semantic reward.

576 Credit-reward factorial follow-up. To separate the
577 two changes in that pilot, we completed a four-arm,
578 group-four comparison from the same SFT initializa-
579 tion and seed, with 256 training rows and 64 opti-
580 mizer steps per arm. The arms cross trajectory-level
581 GRPO versus turn-level leave-one-out updates with
582 terminal-only versus additional training-time geodesic
583 progress. On a new, disjoint 256-episode val-seen
584 screen (50 scenes; complete coverage and zero inference
585 errors), the outcome-only GRPO control succeeds on
586 97 episodes. The combined turn/progress, terminal-
587 only turn, and progress-augmented GRPO arms suc-
588 ceed on 96, 98, and 106, respectively. Their paired
589 SR/SPL changes against control are $(-0.39, -0.14)$,

$(+0.39, +0.68)$, and $(+3.52, +1.51)$ percentage points.
Although progress-augmented GRPO shows a positive
SR difference, its SPL difference misses the frozen joint
 $+2$ -point advancement gate. Descriptive scene-cluster
95% intervals for that arm span zero: $[-2.24, 9.26]$ SR
and $[-4.21, 7.17]$ SPL points. An independent per-
episode recount agrees with the paired analysis after
a documented null-ID schema repair to the verifier;
the reward, evaluator, and metric formulas were un-
changed. These one-seed development results do not
establish unseen-scene benefit, and the turn-level and
trajectory-level rules also differ in effective token up-
date scale. The process signal uses privileged geodesic
distance, not the semantic event verifier.

Token-scale-matched terminal credit. The factorial's
turn-level estimator divided each turn contrast by its
action-token count, whereas the trajectory GRPO con-
trol applied an order-one advantage to each action to-
ken. We therefore tested a new group-four terminal-
outcome-only leave-one-out estimator that standard-
izes the active-peer contrast and broadcasts it to each
action token in the turn. From the same SFT ini-
tialization, it ran 64 steps on 256 train rows. On a
third, disjoint 256-episode val-seen development screen
(38 scenes), candidate/control successes are 108/103;
paired SR and SPL are $+1.95$ and $+1.26$ percent-
age points. Each arm covers every episode once with
no inference errors, and an independent raw recount
agrees. Exploratory scene-cluster intervals span zero:
 $[-4.50, 8.04]$ SR and $[-5.11, 7.26]$ SPL points. Both
improvements fall short of the predeclared joint $+2$ -
point gate, so the original conditional three-seed ex-
pansion was skipped. A subsequent user-directed, ex-
plicitly post-result group-eight sensitivity and group-
four three-seed expansion were conducted separately.
This one-seed screen does not establish unseen-scene
benefit or validate a semantic reward.

Post-result group-size sensitivity. Keeping the 64-
step, 256-row prompts and terminal-RLOO rule fixed,
we sampled eight rather than four trajectories per in-
struction. This doubles simulator rollout work. The
group-eight run had logged nonzero actor gradients on
63 of 64 steps; the remaining step had an all-failure
group and zero advantage. On the same 256-episode
val-seen screen, the group-eight checkpoint succeeds on

112 episodes, versus 103 for the destination-only control and 108 for the group-four checkpoint. Its paired SR/SPL differences are +3.52/+2.82 points versus control and +1.56/+1.56 points versus group four. Every arm covers all 256 episodes across 38 scenes, with no inference errors, and an independent raw-stat recount agrees with the compact episode records. Descriptive scene-cluster intervals for group eight versus control span zero: $[-4.92, 12.13]$ SR and $[-5.28, 11.10]$ SPL points. This comparison was requested after the group-four advancement gate failed; the group-eight result does not retroactively satisfy that gate. It is one-seed, adaptively selected development evidence, with group size confounded by rollout compute.

Post-result group-four scale. The separately authorized expansion used the terminal-RLOO rule with group size four, 512 training rows, and 128 optimizer steps for each of seeds 11, 22, and 33. Every candidate and its same-seed destination-only control covered all 1,839 val-unseen episodes (11 scenes), with zero inference errors. Paired candidate-minus-control SR/SPL differences are +3.97/+3.30, -3.26/-3.07, and +4.13/+3.00 percentage points, respectively. Their mean differences are +1.61/+1.07 points, with seed standard deviations 4.22/3.60 points. Exploratory seed-and-scene intervals span zero: $[-3.03, 5.12]$ for SR and $[-2.86, 4.21]$ for SPL. The compact per-episode recount agrees with the raw-stat verifier. The opposite sign for seed 22 and the intervals preclude a robust-benefit claim. Val-unseen had already influenced method development, so this is larger development replication rather than an untouched test. It does not establish that the group-eight pilot will generalize or scale.

Multimodal GAE actor-critic pilot. We also tested a PPO actor [13] with masked token-level GAE [12], a separately initialized multimodal value head, and terminal-outcome-only feedback. The group-four actor and critic started from the same SFT model; 64 optimizer steps on 256 train rows had finite value losses and nonzero actor and critic gradients at every step. On an adaptively reused 778-episode val-unseen development screen (53 scenes), the GAE candidate succeeds on 22 episodes versus 272 for the same-seed destination-only control: paired SR -32.13 and SPL -31.56 percentage points. Both arms have exact episode coverage and zero inference errors; two independent raw-stat recounts and a compact export recount agree. Descriptive scene-cluster intervals are $[-36.83, -27.43]$ SR and $[-36.19, -26.95]$ SPL points. Mean path lengths are 2.79 versus 6.53 m, consistent

Table 6. Exploratory 128-step policy-prefix follow-up on all 1,839 val-unseen episodes. Differences are branching minus a same-data, from-scratch destination-only control, in percentage points. Mean $\pm$ sample standard deviation is over three training seeds; each model has one stochastic decode per episode.

Seed	Δ SR	Δ SPL
11	-7.18	-6.86
22	+0.49	+0.80
33	+3.26	+4.41
Mean $\pm$ SD	-1.14 ± 5.41	-0.55 ± 5.75

with shorter candidate trajectories but insufficient to identify why learning failed. The prespecified joint +2-point scale gate failed, so the conditional three-seed GAE expansion was skipped. This one-seed development result is not evidence about unseen-scene generalization. The frozen recurrence discounts each generated action-text token ($\gamma = 0.99$, $\lambda = 0.95$), skipping observation spans. A post-result synthetic zero-value check shows text-length-dependent early credit at a fixed decision count. Decision-clock GAE remains untested, and this diagnostic does not identify the cause of the navigation decline.

C. Start-state and rollout-count comparisons

Exploratory policy-prefix follow-up. Separately from EventTrace, we tested whether changing the training-state distribution helps destination-only GRPO. For each of three seeds, a branching curriculum replayed a training-only policy prefix before sampling two continuations; its matched control started at the task origin. Both arms used the same 512 train episodes, SFT initialization, destination reward, and 128 optimizer steps. All six checkpoints covered the complete 1,839-episode val-unseen split with zero inference errors. Table 6 shows sharply different seed effects and a negative three-seed mean. Exploratory scene-and-seed 95% intervals for the mean differences are $[-6.51, 4.02]$ SR and $[-6.16, 4.85]$ SPL points. On the 1,583 episodes outside the fixed 256-episode development screen, the mean changes remain -1.22 SR and -0.55 SPL points. Separate stochastic decodes of the overlapping 256 episodes changed success outcomes for 20-41 episodes per checkpoint, so the small screen is an unstable model-selection signal. This follow-up changes start state and remaining budget together; it does not identify why EventTrace failed or establish a benefit from branching. The checked paired episode records and training audits accompany the manuscript.

Table 7. Exploratory 128-step, destination-only group-size follow-up on all 1,839 val-unseen episodes. Differences are four samples minus a matched two-sample control, in percentage points; the four-sample arm uses twice the simulator rollouts.

Seed	Δ SR	Δ SPL
11	+0.65	+0.96
22	+1.36	+1.44
33	+4.30	+5.12
Mean $\pm$ SD	$+2.10 \pm 1.93$	$+2.51 \pm 2.27$

Group-four optimizer follow-up. A separate destination-only comparison tests four sampled trajectories per GRPO group against the same-data, two-sample control. Both arms use the same SFT initialization, 512 unique training episodes, and 128 optimizer steps at each of three matched seeds; the four-sample arm uses twice as many simulator rollouts. All six checkpoints cover the identical 1,839-episode val-unseen manifest with zero inference errors. Table 7 shows positive paired differences in each seed: the mean is $+2.10 \pm 1.93$ SR and $+2.51 \pm 2.27$ SPL percentage points (mean $\pm$ sample standard deviation over seeds). On the 1,583 episodes outside the fixed 256-episode screen, the means are $+1.94$ SR and $+2.40$ SPL points. The largest seed-33 gain is $+4.30$ SR points, while exploratory 11-scene bootstrap intervals for seeds 11 and 22 cross zero. One stochastic decode per checkpoint and only three training seeds limit the inference. This comparison does not identify a better semantic reward; subsequent reward comparisons must hold group size at four in both arms. Paired records, training audits, and independent recount code accompany the manuscript.

D. Representation-to-reward development

Instruction-grounded reward representation. We also probe a potential reward that combines a causal progress encoder from frozen navigation-policy states with a SigLIP image-instruction encoder adapted on training-scene goal panoramas and group-four successful/failed routes. SigLIP text is padded to a fixed 64 tokens, and offline scoring uses the same one-image request shape as the online service. Each component is normalized by its fit-scene mean absolute margin; a fixed 1:1 mixture is applied without development-set tuning. On 52 development pairs, successful-route ranking rises from 37/52 for the temporal term and 32/52 for the visual term to 43/52 for the mixture; correct-versus-wrong-instruction preference is 39/52. On 38 new episode-disjoint, train-scene pairs, the mixture scores 30/38 route ranks and 34/38 instruction

comparisons. A previously inspected 48-pair scene audit scores 41/48 and 37/48. These selected offline pairs share train scenes or prior audit use, so they provide a screen for an online group-four reward comparison, not evidence of val-unseen navigation gain.

Online fused-reward screen. We added this frozen potential as a bounded terminal bonus to the destination reward and trained a matched group-four policy for 64 steps at seed 11. Its training audit verifies 1,024 reward requests for 1,024 rollouts across 256 matched train episodes. On the fixed 256-episode val-unseen screen, the fused-reward checkpoint achieves 68/256 successes versus 80/256 for the same-data, same-group-size destination-only control: paired SR changes by -4.69 percentage points and SPL by -4.34 points. Both arms cover exactly 256 episodes in 11 scenes with zero inference errors. This single-seed screen does not support scaling this particular reward, despite its offline ranking results. A failure-aware encoder trained on unsuccessful pairs improves a separate 106-pair train-scene ranking audit from 69/106 to 79/106, but direct fusion with the image-instruction score weakens reused instruction-matching probes. These findings do not establish a navigation gain.

Failure-only reward screen. We next applied the new temporal potential only to unsuccessful rollouts, leaving successful rollouts on the original destination reward. A matched seed-11 group-four policy completed 64 steps on the same training episodes; its audit found 854 reward requests for 854 unsuccessful rollouts and no extra bonus for 170 successful rollouts. On the same fixed 256 val-unseen episodes, it achieves 63/256 successes versus the destination-only control's 80/256: paired SR changes by -6.64 points and SPL by -6.05 points. Both models cover 11 scenes without inference errors. The exploratory scene-cluster 95% intervals are $[-11.07, -2.29]$ SR and $[-10.66, -1.47]$ SPL points. This screen fails the predeclared positive SR-and-SPL scale gate; we therefore do not run a three-seed scale-up for this reward. Training-only failed-pair ranking (622/970 in all-failure groups) does not transfer into a navigation benefit. A post-hoc behavior check finds maximum-turn termination on 187 candidate episodes versus 107 control episodes, while mean traveled path length falls from 6.74 to 5.71 m; these differences motivate a stopping-aware reward test but do not identify the cause of the regression. An offline equal mixture with the adapted visual score also fails its train-scene development gate (68/125 failed-pair ranks versus 90/125 for the temporal term), so its audit remains unopened.

814 Stop-aware reward screen. To test whether terminal-
815 mode conditioning mitigates this failure, we gave the
816 frozen temporal bonus only to unsuccessful voluntary
817 STOP trajectories. Successes retained the destination
818 reward, while turn-limit and format failures received
819 no bonus. The matched seed-11 group-four policy com-
820 pleted 64 steps on the same training episodes. Its audit
821 found exactly 328 scorer requests for 328 unsuccessful
822 voluntary stops across 1,024 rollouts, with nonzero ac-
823 tor gradients at every step. On a new, disjoint 256-
824 episode val-unseen screen covering 10 scenes, the can-
825 didate succeeds on 45/256 episodes versus 93/256 for
826 its same-seed destination-only control. Paired SR and
827 SPL changes are -18.75 and -17.43 percentage points;
828 both arms have zero inference errors. The predeclared
829 positive SR-and-SPL gate failed, so this reward re-
830 ceived no three-seed expansion. Post-hoc analysis finds
831 mean path length 4.44 versus 6.93 m and only three
832 candidate episodes with a non-null termination rea-
833 son versus 116 controls, consistent with frequent early
834 stops. On the candidate’s train rollouts, its frozen raw
835 score has AUC 0.486 for distinguishing 18 unsuccessful
836 STOP endpoints within 3.5 m of the goal from 310 far-
837 ther endpoints; only 11 of 183 positive-bonus stops are
838 within 3.5 m. This train-scene diagnostic does not es-
839 tablish causality, but it shows that failure-pair ranking
840 is insufficient calibration for rewarding STOP.

841 Mode-stratified ordinal screen. To remove a direct re-
842 ward offset between unsuccessful STOP and turn-limit
843 trajectories, we next ranked frozen temporal scores
844 only within the same terminal mode of each four-
845 rollout all-failure group; any-success groups retained
846 the original outcome reward. This is a trajectory-
847 level group comparison, not true turn-wise credit as-
848 signment: the current GRPO implementation broad-
849 casts a group-normalized total reward to all action
850 tokens. The 64-step seed-11 training audit verified
851 1,024 rollouts, 149 all-failure groups with active ordi-
852 nal rewards, 827 scorer requests for exactly 827 fail-
853 ures, and a nonzero actor update at every step. On
854 a third fixed, nonoverlapping 256-episode val-unseen
855 screen spanning nine scenes, the mode-stratified can-
856 didate succeeds on 59/256 episodes versus 73/256 for
857 its same-data, same-seed group-four outcome-only con-
858 trol. Paired SR and SPL change by -5.47 and -5.23
859 percentage points, with zero inference errors in either
860 arm. The exploratory nine-scene bootstrap 95% in-
861 tervals are $[-9.06, -1.95]$ SR and $[-8.98, -1.74]$ SPL
862 points. Among failures, the candidate ends with an
863 unforced model STOP 148 times versus 84 for the
864 control; this post-hoc association does not identify a
865 causal mechanism. The predeclared positive SR-and-

SPL scale gate failed, so no 128-step or three-seed ex-
tension was run for this reward. Absolute success rates
across the three 256-episode screens are not directly
comparable because their episode sets differ.

Privileged turn-wise credit upper bound. The preced-
ing trajectory-level rewards broadcast one return to ev-
ery action token. To isolate whether finer credit assign-
ment can help, we tested a training-only oracle that
labels each executed turn by its change in simulator
geodesic distance. In all-failure groups of four, it cen-
ters return-to-go among active rollouts at each turn
and gives no auxiliary advantage to generated STOP
tokens; groups containing a success retain destination-
outcome GRPO. Geodesic distance is never a policy
observation, but this oracle is not a learned seman-
tic reward. The same-data seed-11 pilot completed 64
steps. An independent training audit checked 256 four-
rollout groups, 11,450 turns, active contrast in all 141
all-failure groups, and nonzero actor gradients on all 64
steps. On a fourth, pre-frozen 256-episode val-unseen
screen spanning nine scenes, it succeeds on 81 episodes
versus 72 for its outcome-only control: paired SR $+3.52$
and SPL $+3.76$ percentage points, with zero inference
errors. The exploratory scene bootstrap 95% inter-
vals span $[-2.73, 9.13]$ SR and $[-2.58, 9.28]$ SPL points.
This one-seed screen selected a same-checkpoint full
val-unseen recheck. Across all 1,839 episodes, the ora-
cle succeeds on 602 versus 545 for the control: paired
SR $+3.10$ and SPL $+3.05$ points. On the 1,583 episodes
outside the reused screen, it succeeds on 521 versus
481: SR $+2.53$ and SPL $+2.50$ points. Both evalu-
ations have exact coverage and zero inference errors.
This remains a post-screen, one-seed development re-
sult using privileged training labels; it neither validates
EventTrace’s semantic verifier nor establishes a deploy-
able learned reward.

We next scaled this privileged mechanism to 512
distinct training episodes and 128 optimizer steps at
each of three seeds, with four rollouts per group in
both oracle and destination-only arms. All six check-
points were decoded once on the same 1,839 val-unseen
episodes, with exact coverage and zero inference errors.
Table 8 gives paired results. The mean oracle change is
 $+2.47 \pm 3.18$ SR and $+2.51 \pm 2.78$ SPL points (sample
SD over seeds); seed 22 instead loses 1.09 SR points.
On the 1,583 episodes outside the reused 256-item de-
velopment screen, the means are $+2.63$ SR and $+2.64$
SPL points. Exploratory scene-and-seed intervals for
the full-set mean are $[-1.49, 6.05]$ SR and $[-1.20, 5.78]$
SPL points. Thus the observed mean is positive, but
the seed variation and intervals do not support a con-
sistent gain. The reward uses simulator distance.

Table 8. Matched group-four, 512-row, 128-step privileged turn-wise oracle versus destination-only control on all 1,839 val-unseen episodes. Differences are oracle minus control in percentage points.

Seed	SR counts O/C	Δ SR	Δ SPL
11	513/450	+3.43	+3.45
22	525/545	-1.09	-0.61
33	585/492	+5.06	+4.70
Mean $\pm$ SD	—	$+2.47 \pm 3.18$	$+2.51 \pm 2.78$

An adaptive stopping-boundary pilot tested whether capping geodesic progress credit inside the 3 m success region and penalizing further movement there by 0.1 improves this oracle. Its seed-11 matched $n=4$ 64-step audit verified 256 identical training rows, nonzero actor gradients at every step, and 902 inside-boundary penalty turns. On the reused 256-episode development screen, the candidate succeeds 67 times versus 72 for the unchanged control: paired SR -1.95 and SPL -1.62 points, with zero inference errors. This fails the fixed positive SR-and-SPL scale gate. Neither privileged variant establishes a deployable semantic reward.

Two observation-only progress representations were screened before online use. An instruction-conditioned scalar potential fit to cached policy RGB histories reached only 48.88% balanced forward/regression accuracy on its small train-scene development split. A joint two-view antisymmetric LoRA reached 53.52% balanced accuracy and 52.08% correct-versus-wrong-instruction preference at its selected checkpoint; at 9.88% stationary false positives, its forward recall was zero. Neither passed the fixed development gate or entered RL. A further start-anchored, action-conditioned three-view LoRA also failed its small development gate (55.51% balanced direction, 22.22% regression accuracy), so it received no audit or online reward test. This leaves distillation of the privileged credit signal into an observation-only representation as open work.

We also screened arrival evidence on a scene-disjoint R2R-train split of audited group-four policy rollouts. A single-frame occupancy head reached AUC 0.6053, but recalled only 8.33% of inside-region views and 5.04% of inside-region views on unsuccessful routes at a pooled false-positive rate below 5%. A separate joint two-view, instruction-conditioned transition head was trained for a fixed 1,024 microsteps on real crossings, far nonarrivals, real near-goal retreats, and identical-view wrong-instruction pairs. On all 758 development pairs from 351 records, its selected checkpoint reached AUC 0.6430, versus 0.5177 for a text-only control and 0.5211 after same-scene image-pair shuffling. For identical image pairs with different in-

structions, the correct instruction ranked higher in only 98/176 cases (55.68%). A CPU-only source audit confirms that these pairs share the exact simulator states; their wrong-instruction labels are geometric alternate-goal proxies, without independent semantic adjudication. Weak grounding and label ambiguity cannot be separated by this diagnostic. Nevertheless, at 4.89% pooled false positives it recalled only 15.63% of crossings and 13.45% of unsuccessful crossings, below the frozen 55% and 50% gates. These exploratory screens reused development scenes across method design; neither opened its reserved train-scene audit or entered online RL. Better AUC therefore supplies no evidence of a navigation gain from this learned reward.